

Convolutional Neural Networks Basics

Computer Vision Problems

- Image classification
- Object detection
- Neural STYLE Transfer (Repaint the content image to a style image)



Deep Learning on Large Images

- Challenges of computer vision problems: inputs can get really big
- Ex) worked on $64 \times 64 \times 3$ images - three color channels
- x has input features 12288 \rightarrow but this is a small image
- Ex) $1000 \times 1000 \times 3 = 3$ billion parameters
- Difficult to get enough data to prevent a neural network from overfitting



Edge Detection Example

- Convolution operation
- Early layers of neural network - detect edges
- Later layers - detect partial objects
- Even later layers - complete objects
- 1. Detect vertical edges
- 2. Detect horizontal edges



Vertical Edge Detection

- 6 x 6 x 1 Gray scale image
- Construct 3 x 3 filter (kernel)
- Take 3 x 3 filter and paste it on top of the 3 x 3 region of original input image
- Take filter and shift it to the right

“convolution”

3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

*

1	0	-1
1	0	-1
1	0	-1

=

-5	-4	0	8
-10	-2	2	3
0	-2	-4	-7
-3	-2	-3	-16

3 ¹	0 ⁰	1 ⁻¹	2	7	4
1 ¹	5 ⁰	8 ⁻¹	9	3	1
2 ¹	7 ⁰	2 ⁻¹	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

Add up the 9 numbers = -5

Vertical Edge Detection

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0




*

1	0	-1
1	0	-1
1	0	-1




=

0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0




Vertical Edge Detection Example

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0




*

1	0	-1
1	0	-1
1	0	-1




=

0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0




10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0



*

1	0	-1
1	0	-1
1	0	-1



=

0	-30	-30	0
0	-30	-30	0
0	-30	-30	0
0	-30	-30	0
0	-30	-30	0



Vertical and Horizontal Edge Detection

1	0	-1
1	0	-1
1	0	-1

Vertical

1	1	1
0	0	0
-1	-1	-1

Horizontal


10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10

*

1	1	1
0	0	0
-1	-1	-1

=

0	0	0	0
30	10	-10	-30
30	10	-10	-30
0	0	0	0



Different Kinds of Filters

1	0	-1
1	0	-1
1	0	-1

1	0	-1
2	0	2
1	0	-1

3	0	-3
10	0	-10
3	0	-3

Sobel filter

**Puts more weight to the central pixel
Makes it more robust**

Schorr filter

**For vertical edge detection
(flip 90 degrees for horizontal detection)**

3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

W1	W2	W3
W4	W5	W6
W7	W8	W9

**Don't need to pick, make it parameters
Learn using back propagation**

**Goal: learn 9 parameters
Take image, convolve with 3 x 3 filter
-> good edge detector**

Underlying convolution operation

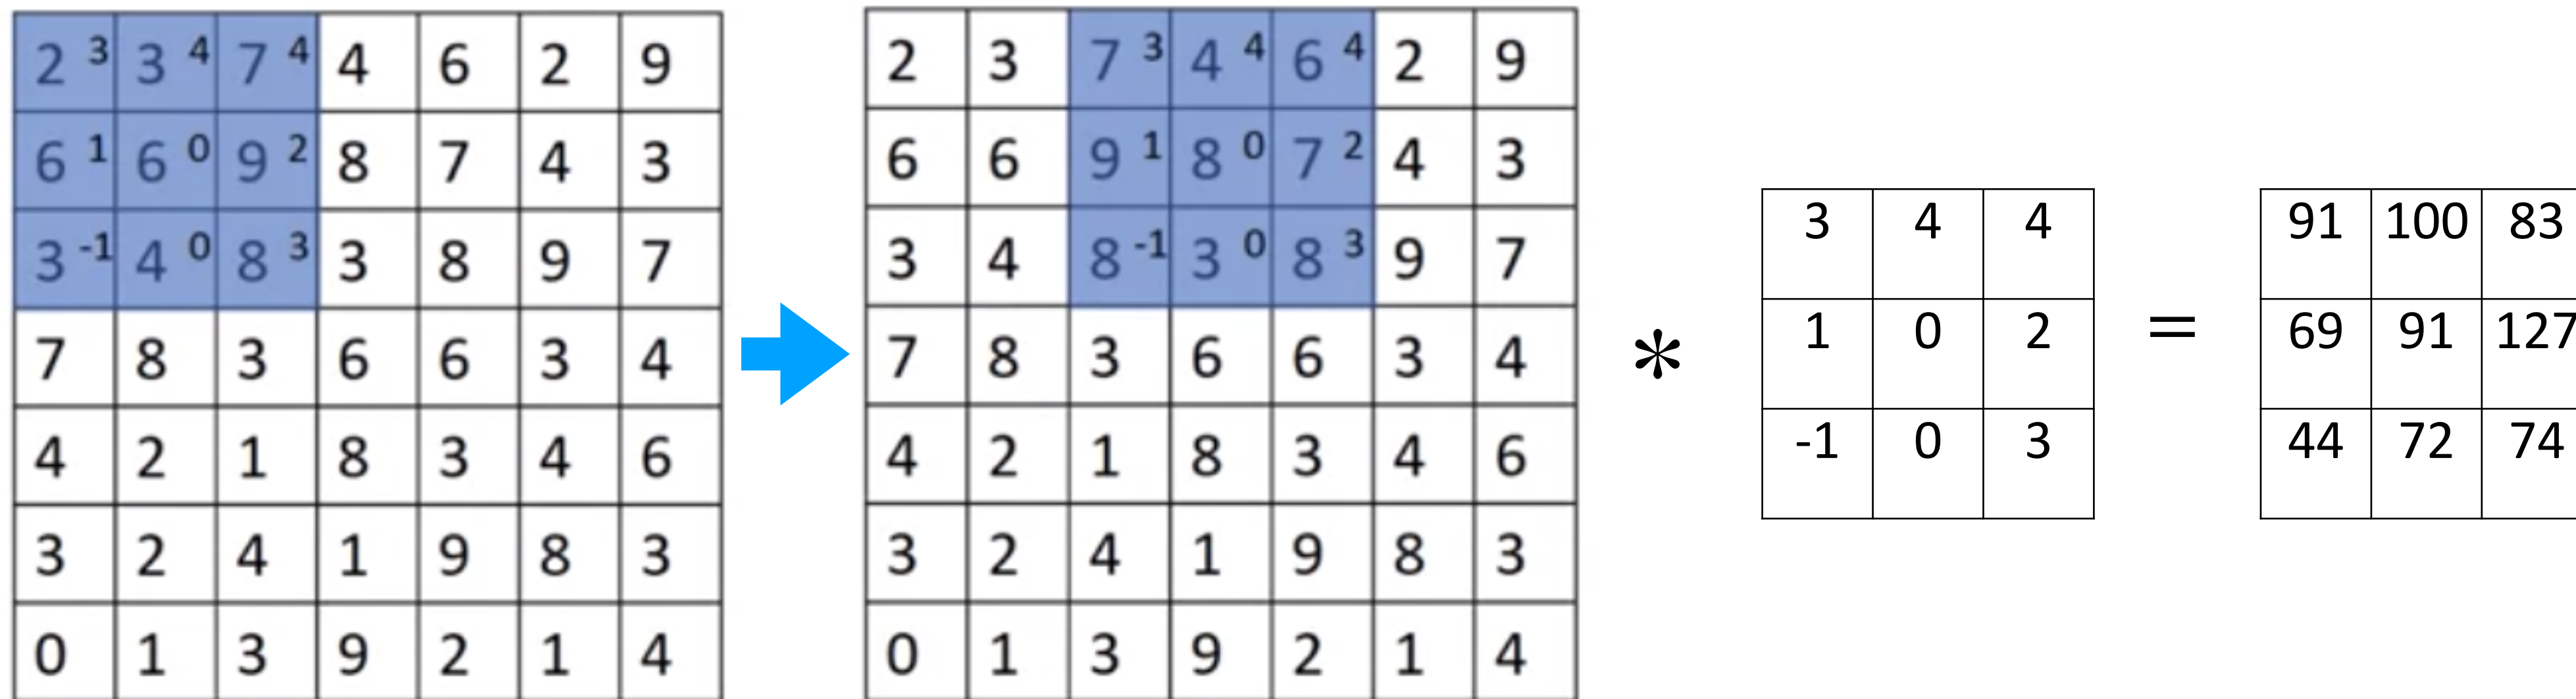
Padding

- Modification to the basic convolutional operation
- $n \times n * f \times f = (n-f+1) \times (n-f+1)$
- Dimension of the output shrinks & counting less information from edges of the image (pixels in corners less in output)
- Solution: before applying convolutional operation, pad the image with an additional one border
- Ex) Pad with zeros, one pixel all around. $P = \text{padding} = 1$
- $(n + 2p - f + 1) \times (n + 2p - f + 1)$

Valid and Same Convolutions

- “Valid”: no padding. $n \times n * f \times f = (n-f+1) \times (n-f+1)$
- “same”: pad so that output size is the **same** as the input size.
 - $(n+2p-f+1) \times (n+2p-f+1)$
 - $P = (f-1)/2$
 - F is usually odd: it has central pixel, symmetric padding

Strided Convolutions



Stride = 2

Output size

If $(n + 2p - f)/s + 1$ not an integer, round it down.
 The filter must lie entirely within the image (+padding)

$n \times n * f \times f = (n + 2p - f)/s + 1 + (n + 2p - f)/s + 1$
 Image Filter Padding Stride